

Heuristic Optimization: implementation exercise 1

Samuel Buchet: 000447808

March 2017

1 Iterative improvement heuristic

1.1 Implementation details

The implementation has been done in c++ from the sample code. One improvement function has been written for each neighborhood. The improvement functions take a parameter that decides the pivoting rule (first or best improvement). For the first improvement rule, after a new improvement, the algorithm doesn't restart the search from the first neighbor, as seen in class.

One of the main difficulty of the implmentation was the computation of the cost of a solution (weighted completion times sum). The easiest way is to compute the entire sum, starting from the first job. However, it takes a lot of time and it is not always necessary. For example, if only the last two jobs are modified in the solution, only the last two completion times need to be re-computed. To recompute only a part of the solution cost, it is necessary to keep the completion times of each job for each machine. These times are saved in a matrix in the *PfspInstance* class. Thanks to this matrix, two functions of the class *PfspInstance* are available to compute the completion time, one that computes the the total weighted sum and modify the matrix, and another that computes the sum without modifying the matrix.

These two functions are used in the iterative improvement procedures. The iterative improvement functions have been written for each neighborhood. The procedures iterate over all neighbours. If the pivoting rule is the best improvement rule, the best neighbor is saved and the change is done at the end of the procedure. If the pivoting rule is the first improvement rule, each time an improving neighbor is found, the change is done and the search continue with the new solution.

The transpose improvement procedure is the easiest in the sense that there is no nested loop. In this procedure, all the possible transpose moves are tried on the solution. As explained below, it is not necessary to recompute the wighted sum from the first job if the transpose move is not done at the begining of the solution. However, to do so, the matrix of the completion times of the *initial* solution has to be kept. That's why the procedure used to compute the weighted sum of a neighbor doesn't modify the matrix. If the pivoting rule is the first improvement rule and the neighbor improve the solution, the cost is re-computed in order to apply the changes on the completion times matrix (with a second procedure which modify the completion times matrix).

The exchange improvement procedure is very similar. The only difference is the nested loop. It is then obvious that this neighborhood is longer to compute than the transpose one. The insert procedure is also similar. However, the insertion is properly done only once. For a given element, the procedure insert it in the first position and the other positions are obtained with transpose moves, which is equivalent. However, the procedure tries the original position of the element and nothing has been done to prevent this case.

All these improvement procedures return a boolean that indicates whether an improving solution has been found or not. This boolean is used in the iterative improvement procedure to stop

the algorithm. The iterative improvement procedure also initialize the solution with the random permutation or the simplified rz heuristic, depending on the parameters of the algorithm.

1.2 Experiments

The following table reports the average of the computation times and the average of the percentage deviations from the best known solution for each algorithm. best. imp means that the pivoting rule is the best improvement rule and first imp. corresponds to the first improvement rule. random means that the solution is initialized with a random permutation and srz means that the simplified rz heuristic has been used. The last parameter is the neighborhood used in the algorithm.

algorithm	average computation time	average percentage deviation
best imp./random/exchange	919.706	4.383
best imp./random/insert	2721.295	3.242
best imp./random/transpose	12.753	36.135
best imp./srz/exchange	239.501	3.157
best imp./srz/insert	872.557	2.231
best imp./srz/transpose	7.757	4.304
first imp./random/exchange	93.031	3.252
first imp./random/insert	292.007	2.676
first imp./random/transpose	1.963	36.163
first imp./srz/exchange	80.713	3.153
first imp./srz/insert	215.415	2.185
first imp./srz/transpose	7.216	4.306

We can see that some algorithms seem to be much worst. Indeed, the average percentage deviation is quite high for the algorithms using random and transpose parameters. We can also notice that the algorithms using the transpose neighborhood are much faster. The algorithms with the insert neighborhood are also the slowest.

1.3 Statistical tests

The following statistical tests have been done using the software R. The statistical test used is the Wilcoxon test with a significance level equal to 0.05.

1.3.1 Initial solution

The following table shows the result of the test applied to all pairs of algorithms which differ only between the initial solution.

algorithms compared	p-value	null hypothesis
best imp. ; exchange	6.173616e-09	rejected
best imp. ; insert	6.35434e-07	rejected
best imp. ; transpose	1.671329e-11	rejected
first imp. ; exchange	0.5339054	not rejected
first imp. ; insert	0.0005783928	rejected
first imp. ; transpose	1.671329e-11	rejected

For most of the tests, the null hypothesis is rejected, which means that there is a statistically significant difference between the solution quality. Moreover, with the srz heuristic, the average of the percentage deviation is lower for each algorithm. We can then conclude that the simplified rz heuristic gives better results than the random permutation.

1.3.2 Pivoting rule

In this section, the test has been applied to all pairs of algorithms which differ only between the pivoting rule. This first table contains the results of the tests applied to the average of the percentage deviation.

algorithms compared	p-value	null hypothesis
random ; exchange	6.147428e-08	rejected
random ; insert	0.001416562	rejected
random ; transpose	0.7100708	not rejected
srz ; exchange	0.6481444	not rejected
srz ; insert	0.463876	not rejected
srz ; transpose	0.9878667	not rejected

We can see that there is a statistically significant difference only with two algorithms (which don't use the best initial solution). Moreover, for these two pairs of algorithms, the first improvement rule seems to give better results.

The table below contains the result of the statistical tests applied to the computation times.

algorithms compared	p-value	null hypothesis
random ; exchange	1.671329e-11	rejected
random ; insert	1.671329e-11	rejected
random ; transpose	1.671329e-11	rejected
srz ; exchange	5.278219e-11	rejected
srz ; insert	1.671329e-11	rejected
srz ; transpose	3.163469e-06	rejected

For each pair of algorithms, there is a statistically significant difference between the computation times. The average of the computation times is lower for the first improvement algorithm for each pairs. We can conclude that the first improvement algorithms are faster and there is no big differences between the quality of the solutions.

1.3.3 Neighborhood

In this section, the different neighborhoods are compared. Comparisons are done for each pair of neighborhood (transpose/exchange, transpose/insert, exchange/insert). The p-value returned by the tests of the computation times and the solution quality are indicated in the tables below.

Comparison between the transpose neighborhood and the exchange neighborhood :

algorithms compared	relative percentage deviation	computation time
first imp. ; random	1.671329e-11	1.671329e-11
first imp. ; srz	5.282725e-11	1.671329e-11
best imp. ; random	1.671329e-11	1.671329e-11
best imp. ; srz	1.138656e-10	1.671329e-11

Comparison between the transpose neighborhood and the insert neighborhood :

algorithms compared	relative percentage deviation	computation time
first imp. ; random	1.671329e-11	1.671329e-11
first imp. ; srz	1.671329e-11	1.671329e-11
best imp. ; random	1.671329e-11	1.671329e-11
best imp. ; srz	1.671329e-11	1.671329e-11

Comparison between the exchange and the insert neighborhood :

algorithms compared	relative percentage deviation	computation time
first imp. ; random	0.0006803447	1.671329e-11
first imp. ; srz	2.115717e-09	1.671329e-11
best imp. ; random	1.345635e-06	1.671329e-11
best imp. ; srz	7.033074e-10	1.671329e-11

We can see that for each comparison, the null hypothesis is rejected. There is a statistically significant difference between the neighborhood in terms of quality and computation time.

For each pair of algorithm, we can see on the table in 2.2 that in average the transpose neighborhood is faster but the quality of the solution found by the exchange neighborhood is better.

If we compare the transpose and the insert neighborhoods, we can see that once again, the transpose neighborhood is faster but its quality is worst than the insert neighborhood in average.

The comparison of average values between the exchange neighborhood and the insert neighborhood indicates that the insert neighborhood is slower than the exchange neighborhood but the quality of its solution is better.

To conclude, the insert neighborhood is the one that generates the best solutions in terms of quality. However, the transpose neighborhood is the fastest neighborhood. The exchange neighborhood is a good tradeoff because it is faster than the insert one and its quality is better than the transpose one.

2 Variable Neighborhood Descent heuristic

2.1 Implementation details

The implementation of the VND algorithm mainly uses the procedures created for the iterative improvement algorithms. The algorithm has been implemented in the *VNDHeuristic* procedure. It first initialize the solution with the chosen initialization algorithm. Then, it applies successively the different iterative improvement procedures following the chosen order of the neighborhoods. This step is repeated until there is no improvement in the last neighborhood.

2.2 Experiments

The table below contains the average of the computation time, the average of the percentage deviation of the solution, the average of the percentage improvement over the exchange neighborhood and the average of the percentage improvement over the insert neighborhood. The order used for the neighborhood is transpose - exchange - insert and transpose - insert - exchange. ex. first means that the exchange neighborhood use applied before the insert neighborhood.

algorithm	average computation time	avg percentage deviation	avg imp exchange	avg imp. insert
random ; ex. first	382.526	2.634	0.006	0.0003
random ; ins. first	307.123	2.091	0.011	0.006
srz ; exchange first	275.016	2.129	0.010	0.0005
srz ; insert first	283.463	2.161	0.010	0.0002

2.3 Statistical tests

The table below contains the p-values obtained with a Wilcoxon test on the vnd algorithms for the quality of the solution and the computation time.

algorithm	computation time	percentage deviation
srz	0.1720717	0.6274306
random	0.000156713	0.0001869606

The tests indicate that there is no statistically significant differences between the solution qualities and the computation times for the srz heuristic. However, for the random permutation as initial solution, there is a difference. For the random permutation, the average of the computation time and the average of the percentage deviation are lower with the insert neighborhood before the exchange neighborhood. We can conclude that this order is better. However, regarding the comparison between vnd and the iterative improvement algorithms, the improvement of vnd seems to be very low.

3 Compilation and execution

To compile the project, make sure you are in the *code* folder and type the *make* command.

To execute the program, type `./flowshopWCT` in the terminal. You need to specify a valid instance. You can also specify options to choose the neighborhood, the pivoting rule and the initialization of the solution. For example, the command `./flowshopWCT yourInstance --srz --first --insert` executes an iterative improvement algorithm with the simplified rz heuristic as initial solution, the first improvement as pivoting rule and the insert neighborhood. Here are all the options available :

- neighborhood : `--insert`, `--exchange`, `--transpose`
- initial solution : `--srz`, `--random`
- pivoting rule : `--first`, `--best`

You can also execute a vnd algorithm. To do that, use the `--vnd` option. To choose the order of the neighborhood, you can precise `--exfirst` and `--insfirst`. And you can still indicate which initial solution to use (vnd uses first improvement rule).